

Profiling Ruby

with Flamegraphs

Janusz Mordarski

\$ whoami

— — —

Janusz Mordarski

#ruby #rails #python #javascript #go #web #api #devops #rest
#amqp #unix #agile #kanban #scrum

janusz@mordarski.eu

<https://github.com/januszm>

<https://www.linkedin.com/in/janusz-mordarski/>

<http://mordarski.eu>

Code profiling

Why?

— — —

“if you can’t **measure** it, you can’t **improve** it”

“**know** your system”

Why is my code slow?

(when should I ask this question?)

What?

— — —

What is profile ?

“To **profile** someone means to give an account of that person's life and character.”

“Your **profile** is the outline of your face as it is seen when someone is looking at you from the side.”

“A record of a person's psychological or **behavioural characteristics**, preferences, etc.”

Source: <https://www.collinsdictionary.com/dictionary/english/profile>



R. Lewandowski (ID: 188545)

Robert Lewandowski **ST** Age 28 (Aug 21, 1988) 6'1" 174lbsOverall Rating **91**Potential **91**

Value €92M

Wage €355K

Preferred Foot **Right**International Reputation **4** ★Weak Foot **4** ★Skill Moves **3** ★Work Rate **High / Medium**Body Type **Normal**Real Face **Yes**Release clause **€151.8M** **FC Bayern Munich****86** ★★★★★Position **ST**Jersey Number **9**Joined **Jul 1, 2014**Contract Valid Until **2021** **Poland****78** ★★★★★☆Position **ST**Jersey Number **9**

#Clinical Finisher

Follow (560)

Like (0)

Dislike (16)

Jump to

Attacking

62 Crossing**91** Finishing**85** Heading Accuracy**85** Short Passing**89** Volleys

Skill

85 Dribbling**77** Curve**84** FK Accuracy**65** Long Passing**89** Ball Control

Movement

79 Acceleration**79** Sprint Speed**78** Agility**91** Reactions**80** Balance

Power

88 Shot Power**84** Jumping**79** Stamina**84** Strength**84** Long Shots

What?

— — —

What can we profile?

[https://en.wikipedia.org/wiki/Profiling_\(computer_programming\)](https://en.wikipedia.org/wiki/Profiling_(computer_programming))

What matters for “internet” applications?

- **CPU** usage
- memory consumption
- what else?


Call stack

— — —

Who has the control over the CPU?

```
net => OS => http server =>  
def call  
  @record = Model.find_by(uuid: ...)  
  begin  
    crm.set_paid(@record.op_id, @transaction[:id])  
  rescue StandardError => e  
    Raven.capture_exception(e, extra: { id: ..., ...  
  end  
  AgentMailer.new_application(@household).deliver_later  
end
```



http server => OS => net



Two approaches (that matter) [for us]

Profilers:

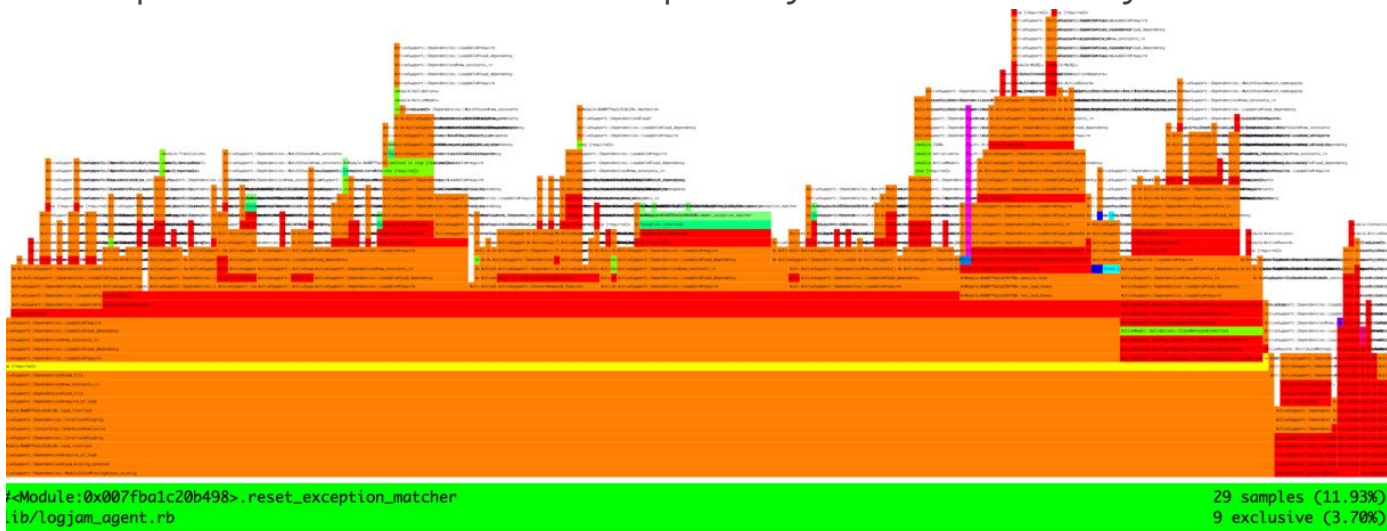
- **instrumenting** / tracing / deterministic
 - from function to function, track the execution time
- **sampling** / statistical
 - from time slice to time slice, call stack snapshots



Seeing is believing

<http://www.brendangregg.com/flamegraphs.html>

“Flame graphs are a visualization of profiled software, allowing the most frequent code-paths to be identified quickly and accurately.”



`is_a? Ruby`

Stackprof

Sampling profiler. Very **low footprint**, low accuracy.

Great for long running processes or methods (eg. processing data in a loop).

<https://github.com/tmm1/stackprof>

ruby-prof

— — —

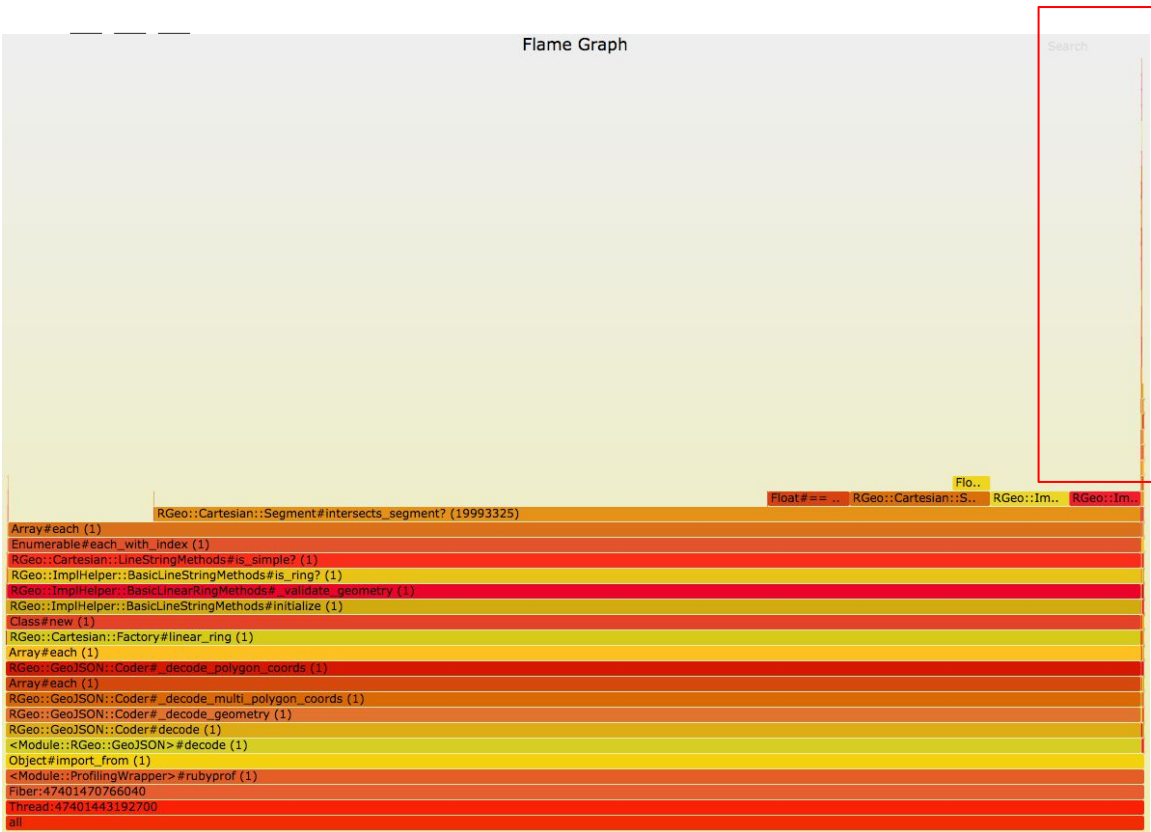
Instrumenting profiler, **accurate**, but may affect the processing time (observer effect, “heisenbugs”).

Great for **pinpointing** issues and profiling short processes or methods.

<https://github.com/ruby-prof/ruby-prof>

sampling vs instrumenting (very short process)






sampling vs instrumenting (short process looped)



																				rgeo	29472x 99.33%
																				ruby-2.4.3	177x 0.60%
																				activesupport	147x 0.50%
																				activerecord	121x 0.41%
																				current	94x 0.32%
																				arel	9x 0.03%
																				ruby	8x 0.03%
																				activemodel	7x 0.02%
																				concurrent	6x 0.02%
																				pg	1x 0.00%

something is wrong here, 99% of the time is spent in RGeo (<https://github.com/rgeo/rgeo>)


sampling vs instrumenting (short process looped)



ruby-2.4.3	280x
rgeo	531.82%
active_record	281x
activesupport	53.68%
current	151x
concurrent	48.27%
ruby	145x
activemodel	84x
arel	22.48%
pg	8x

30k samples
vs
200 samples!

150x
improvement



change: install geos-devel

Configure Elastic Beanstalk
containers.

```
.ebextensions/...
packages:
  yum:
    git: []
    geos-devel: []
    httpd: []
```

Wrapping it up

— — —

```
require "ruby-prof"
require "stackprof"

module ProfilerWrapper
  def self.rubyprof(name)
    GC.disable

    result = ::RubyProf.profile do
      yield
    end

    printer = ::RubyProf::FlameGraphPrinter.new(result)
    printer.setup_options(print_file: true)
    File.open("tmp/#{Time.now.to_i}_#{name}-rubyprof.dump", "w") do |file|
      printer.print(file, {})
    end
  end
end

GC.enable
end
```

```
# Gemfile
gem "ruby-prof"
gem "ruby-prof-flamegraph"
gem "stackprof"
```

Wrapping it up

— — —

```
module ProfilerWrapper
  ...
  def self.stackprof(name)
    GC.disable
    ::StackProf.run(
      mode: :cpu,
      raw: true, # for the flamegraph
      out: "tmp/#{Time.now.to_i}_#{name}-stackprof.dump", # Thread.current.object_id
      interval: 1000 # microseconds
    ) do
      yield
    end
  ensure
    GC.enable
  end
end
```

Wrapping it up

— — —

```
ProfilerWrapper.rubyprof("some_module") do
  ...
end

messages.each do |amqp_message|
  if amqp_message.payload["foo"] == "bar"
    ProfilerWrapper.rubyprof("bar") { process(amqp_message) }
  else
    process(amqp_message)
  end
end
```

```
ProfilerWrapper.stackprof("some_other_module") do
  ...
end

ProfilerWrapper.stackprof("messages") do
  messages.each do |amqp_message|
    process_message(amqp_message)
  end
end
```

Flamegraphs


— — —

```
# RubyProf
$ bin/flamegraph.pl --countname=ms tmp/12193520_1490104692_myname-rubyprof.dump >
tmp/12193520_1490104692_planned-rubyprof.svg

# Stackprof
$ stackprof --flamegraph tmp/11464440_1490025005-stackprof-cpu-myname.dump >
tmp/11464440_1490025005-stackprof-cpu-process_planned.flame
$ stackprof --flamegraph-viewer=tmp/11464440_1490025005-stackprof-cpu-myname.flame
> open
file:///path_to/ruby-2.4.3@mygemset/gems/stackprof-0.2.11/lib/stackprof/flamegraph/viewer.html?data=path_to_project/tmp/11464440_1490
025005-stackprof-cpu-myname.flame
```

(stackprof --flamegraph internally uses the flamegraph.pl script anyway)

[illegible]



Use case (amqp processing, background jobs)

“Planned”, daily push notifications

processing large numbers of records in an event driven environment.

Stack: Rails + AMQP + MySQL


Cause: major increase in the volume of the content pushed to users (push notifications).

Goal: to be 10x faster.

Use case (amqp processing, background jobs)

Steps:

- wrap the important part of the code in the profilers
- collect results (production env, 1 of XX worker servers)
- generate flamegraphs
- refactor the code





activerecord	1932x
	100.00%
models	1892x
	97.93%
activesupport	1884x
	97.52%
mysql2	1040x
	53.83%
ruby-2.3.1	644x
	33.33%
time_bandits	618x
	31.99%
logjam_agent	610x
	31.57%
activemodel	88x
	4.55%
xing	24x
	1.24%
concurrent	14x
	0.72%
arel	12x
	0.62%
tzinfo	4x
	0.21%
whenever	3x
	0.16%

ActiveRecord => SQL


— — —

```
- update_attribute(:enqueued_at, Time.now)
  AMQP::Messenger.publish(:xxx, { foo: :bar }.to_json)

+ update_column(:enqueued_at, Time.now)
  AMQP::Messenger.publish(:xxx, { foo: :bar }.to_json)
```





after:



ActiveRecord in 880 samples

before:



ActiveRecord in 1932 samples

ActiveRecord => SQL

— — —

```
sql = %{ UPDATE ... #{ActiveRecord::Base.sanitize(Time.now) ...  
+ Model.connection.execute(sql)  
  AMQP::Messenger.publish(...
```

multiple AR=>SQL == low memory footprint => batch processing (~50x speed improvement)

Use case (request-response cycle)

```
# Gemfile
gem "rack-mini-profiler"
gem "flamegraph"
gem "stackprof"
```

st?pp=flame

```
before_action do
  Rack::MiniProfiler.authorize_request
end
# OR
around_action :wrap_in_profiler

private

def wrap_in_profiler
  ProfilingWrapper.stackprof("request") do
    yield
  end
end
```

unless Ruby

Use case (Python)

— — —

running ML models with Flask

Inspired by: <http://www.alexandrejoseph.com/blog/2015-12-17-profiling-werkzeug-flask-app.html>

Tools: werkzeug (WSGI) + cProfile / callgrind / qcachegrind

```
+++ b/app/app.py
import pandas as pd
from io import BytesIO
+from werkzeug.contrib.profiler import
ProfilerMiddleware

app = Flask(__name__)
+app.wsgi_app = ProfilerMiddleware(app.wsgi_app,
profile_dir='profile')

@app.route('/', methods=['GET', 'POST'])
def main():
```

Use case (Python)

100.00	0.00	(0)	runapp
100.00	0.00	1	wsgi_app
100.00	0.00	1	full_dispatch_request
100.00	0.00	1	dispatch_request
100.00	0.00	1	main
88.11	0.05	1	expected_rent
88.05	0.37	114	run_model
77.46	24.85	114	get_property_info
56.38	56.38	801	<method 'execute' of 'sqlite...
11.85	0.03	2	expenses
11.32	6.63	456	get_current_stats
6.04	0.00	1	total_dicts
5.97	0.25	228	hoa

77.46	99 546 535	114	run_model (expected_rent.py)
-------	------------	-----	------------------------------

OUGH!

ns	ns per call	Count	Callee
49.53	63 656 140	114	<method 'execute' of 'sqlite3.Cursor'...
3.01	3 870 833	114	get_current_stats (expected_rent.py)
0.07	84 131	114	built-in method 'execute' of 'sqlite3.Connection'...

Use case (Javascript)

Live example:

Chrome Developer Tools >
Performance recording
More Tools > Javascript Profiler

Node.js: <https://nodejs.org/en/blog/uncategorized/profiling-node-js/>

Bonus: APM and Benchmarking

Benchmarking

— — —

```
time = Benchmark.realtime { expensive_method }
```

```
> foo: 0.00286
```

```
> bar: 0.05678
```

```
> save: 1.0366e-05
```

Not sexy

Stats aggregation

```
# gem 'statsd-ruby', require: 'statsd'

$statsd = Statsd.new(STATSD_HOST, STATSD_PORT)
$statsd.namespace = "com.mycompany.myproduct_#{Rails.env}"

$statsd.timing(
  "mymodule.#{something}.expensive_method",
  (time*1000).round
)
```

Graphing FTW!

— — —

Hosted solutions:


NewRelic, DataDog, Logjam (<https://demo.logjam.io>)

Custom builds:


timers from Benchmark => [aggregator] => [graph]

- StatsD/Graphite => Grafana
- InfluxData/InfluxDB => Grafana


Import timings production




Calc timings production



Import timings staging



Calc timings staging



Final thoughts

— — —

When?

Rules of thumb:

- run in production (`!= live`)
- start with `stackprof`, adjust interval
- wrap more code or switch to `ruby-prof` if `stackprof` is accurate enough
- use `rack-miniprofiler` for http requests with html responses
- use `stackprof/ruby-prof` for `{background,worker,scheduled}` jobs
- disable GC

Thank you!

— — —

Questions?

Janusz Mordarski

janusz@mordarski.eu

<https://github.com/januszm>

<https://www.linkedin.com/in/janusz-mordarski/>

<http://mordarski.eu>